



LIBPMEMOBJ-CPP WORKSHOPS

Speaker: Szymon Romik (Intel Data Center Group)

<szymon.romik@intel.com>

May, 2019

Agenda

Remind API

Persistent Memory Programming

- Warmup example (persistent counter)
- Finding bugs related to persistent memory programming
- Converting volatile queue to persistent one
- Hashmap example
- Processing data on persistent memory using map reduce

Using libpmemobj-cpp

Introduction and documentation:

- http://pmem.io/pmdk/cpp_obj/

C++ containers

- <http://pmem.io/2018/11/02/cpp-array.html>
- <http://pmem.io/2019/02/20/cpp-vector.html>
- More containers under development

Libpmemobj manpages:

- <http://pmem.io/pmdk/manpages/linux/master/libpmemobj/libpmemobj.7.html>

libpmemobj – what you will need?

PMEMobjpool *pmemobj_open(const char *path, const char *layout);

void pmemobj_close(**PMEMobjpool** *pop);

PMEMoid pmemobj_root(**PMEMobjpool** *pop, **size_t** size);

int pmemobj_tx_add_range(**PMEMoid** oid, **uint64_t** off, **size_t** size);

int pmemobj_tx_add_range_direct(const **void** *ptr, **size_t** size);

PMEMoid pmemobj_tx_alloc(**size_t** size, **uint64_t** type_num);

int pmemobj_tx_free(**PMEMoid** oid);

void *pmemobj_direct(**PMEMoid** oid);

TX_BEGIN(**PMEMobjpool** *pop) / TX_END

OID_NULL, OID_IS_NULL(**PMEMoid** oid)

libpmemobj-cpp – what you will need?

```
pool<T> pool<T>::open(const std::string &path, const std::string &layout)
```

```
peristent_ptr<T> pool<T>::root()
```

```
void transaction::run(pool_base& pool, std::function<void()> tx, ...)
```

```
peristent_ptr<T> pmem::obj::make_persistent<T>(Args &&... args)
```

```
void pmem::obj::delete_persistent<T>(peristent_ptr<T> &ptr)
```

```
Types: p<T>, peristent_ptr<T>
```

Getting started

Get workshops repo:

```
$ git clone http://github.com/pmemhackathon/2019-05  
$ cd
```

How to compile examples?

- Simply run:

```
$ make
```

- Export your user login in order not to modify commands (you can just copy them)

```
$ export USER=pmdkuser{1..100}
```

Persistent counter – what you should do

Change warmup.cpp to print bigger number every time you run it

- Add variable (a counter) to root struct
- Increment the variable inside „inc” method
- Return new value

Expected result:

```
$ pmempool create obj --layout=warmup -s 100M /mnt/pmem-fsdx0/${USER}/warmup
$ ./warmup /mnt/pmem-fsdx0/${USER}/warmup
1
$ ./warmup /mnt/pmem-fsdx0/${USER}/warmup
2
```

Pmemcheck – persistent memory error detector

Checks for non-persistent stores

Checks for overwrites

Checks for stores made outside of a transaction

Checks for snapshotting the same object in two different threads

Can be found here: <https://github.com/pmem/valgrind>

Pmemcheck – installation and usage

Installation (Already installed on your machine)

```
$ git clone https://github.com/pmem/valgrind  
$ cd valgrind  
$ ./autogen.sh  
$ ./configure [--prefix=/where/to/install]  
$ make install
```

Usage

```
$ valgrind --tool=pmemcheck [valgrind options] <your_app> [your_app options]
```

Find bugs – what you should do

Run:

```
$ pmempool create obj --layout=find_bugs -s 100M /mnt/pmem-fsdax0/${USER}/find_bugs  
$ valgrind --tool=pmemcheck ./find_bugs /mnt/pmem-fsdax0/${USER}/find_bugs
```

Fix bugs reported by valgrind and run valgrind again

Queue – what you should do

Implement a persistent version of queue

It should be based on volatile queue (modify queue.cpp file)

Usage for volatile version:

```
$ ./queue  
$ push 1  
$ pop
```

Usage for persistent version

```
$ pmempool create obj --layout=queue -s 100M /mnt/pmem-fsdx0/${USER}/queue  
$ ./queue /mnt/pmem-fsdx0/${USER}/queue  
$ push 1  
$ show
```

Converting queue step-by-step

1. Open a pool using a **path** variable and supply „queue” as layout.
2. Obtain pointer to the root object
3. Change volatile pointers to persistent ones
4. Change memory allocations/deletions
5. Add transactions

Hashmap

Implement a hashmap with following interface:

- `at(key)`
- `Insert(key, value)`. If the key already exists, overwrite old value with new one
- Assume that hashmap user is calling insert method inside transaction

Do following:

- Use `std::hash` specialization for `pmem::obj::experimental::string` (which is already defined)
- Store buckets for a hashmap in `pmem::obj::experimental::array` container
- For buckets you can use `pmem::obj::experimental::vector` with elements containing Key and Value

Hashmap

To check if it works, compile and run:

```
$ ./simplekv_simple path_to_pool
```

Data oriented design

The approach is to focus on the data layout, separating and sorting fields according to when they are needed.

In general – allows for better utilization of CPU cache

For persistent memory – allows for optimized snapshotting

Try to optimize your hashmap layout

Hashmap

Modify your hashmap to make insert() method atomic for user (use transaction inside insert implementation)

- libpmemobj function:

`PMEMoid pmemobj_pool_by_ptr(const void* addr)`

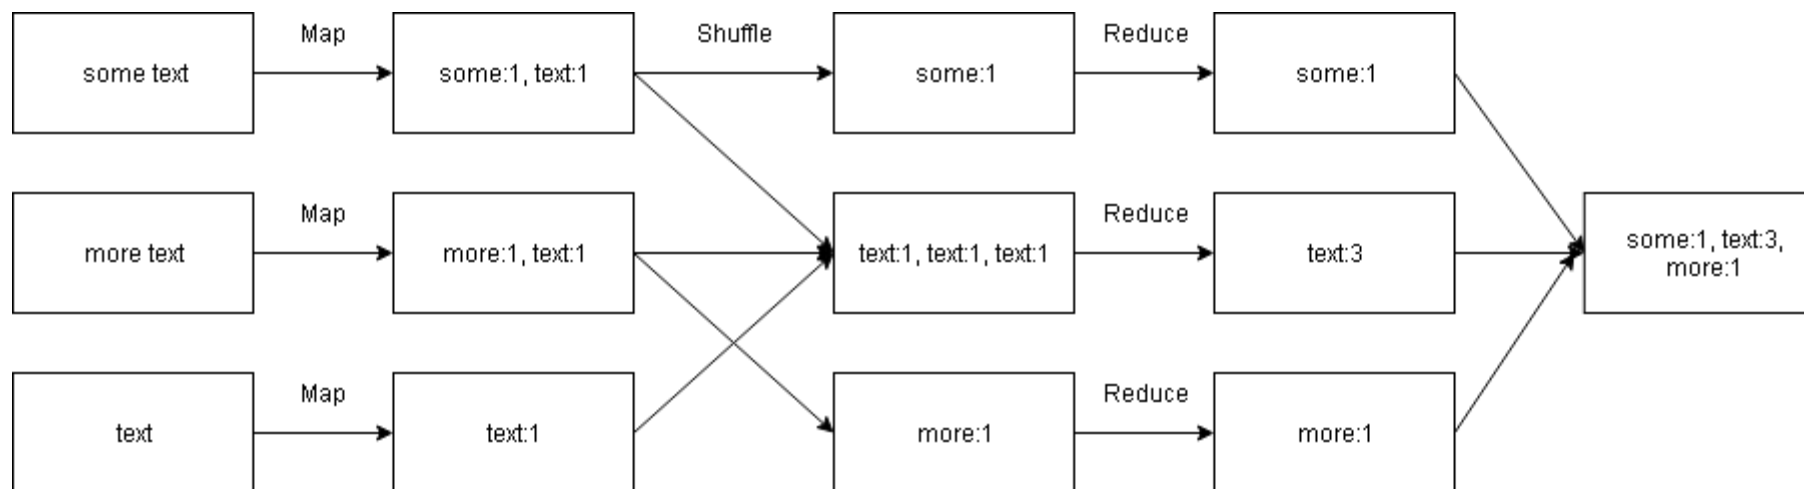
might be useful

Map reduce

Programming model for processing and generating big data sets

Consists of Map, Reduce and Shuffle steps

- Map – performs filtering, transformation or sorting
- Shuffle – redistributes data based on the output keys produced by map step
- Reduce – summary operation (reducing list of values)



Map reduce example

This example uses MapReduce to count words in text files

MapReduce is implemented using:

- `std::transform` - <https://en.cppreference.com/w/cpp/algorithm/transform>
- `std::accumulate` - <https://en.cppreference.com/w/cpp/algorithm/accumulate>

usage (also in README.txt):

```
$ simplekv_word_count pool file1.txt file2.txt ...
```

PMDK Resources

PMDK Resources:

- Home: <https://pmem.io>
- PMDK: <https://pmem.io/pmdk>
- PMDK Source Code : <https://github.com/pmem/PMDK>
- Google Group: <https://groups.google.com/forum/#!forum/pmem>
- Intel Developer Zone: <https://software.intel.com/persistent-memory>
- libpmemobj-cpp: <https://github.com/pmem/libpmemobj-cpp>
- valgrind: <https://github.com/pmem/valgrind>
- NDCTL: <https://pmem.io/ndctl>
- Getting Started Guides: <https://docs.pmem.io>

More Developer Resources

Getting Started

- Intel IDZ persistent memory- <https://software.intel.com/en-us/persistent-memory>
- Entry into overall architecture - <http://pmem.io/2014/08/27/crawl-walk-run.html>
- Emulate persistent memory - <http://pmem.io/2016/02/22/pm-emulation.html>

Linux Resources

- Linux Community Pmem Wiki - <https://nvdimm.wiki.kernel.org/>
- Pmem enabling in SUSE Linux Enterprise 12 SP2 - <https://www.suse.com/communities/blog/nvdimm-enabling-suse-linux-enterprise-12-service-pack-2/>

Windows Resources

- Using Byte-Addressable Storage in Windows Server 2016 - <https://channel9.msdn.com/Events/Build/2016/P470>
- Accelerating SQL Server 2016 using Pmem - <https://channel9.msdn.com/Shows/Data-Exposed/SQL-Server-2016-and-Windows-Server-2016-SCM--FAST>

Other Resources

- SNIA Persistent Memory Summit 2018 - <https://www.snia.org/pm-summit>
- Intel manageability tools for Pmem - <https://01.org/ixpdimm-sw/>

